DHURANDHAR

Project Name : Duality AI’s Offroad Sematic Scene Segementation

PROBLEM STATEMENT

The Challenge

autonomous vehicles in off-road environments, unstructured terrain, similar textures, lighting conditions, and CNN limitations

Objective

building robust semantic segmentation model.

Dataset Structure

Organized into train and val folders for model training and validation.

Image Format

All images are resized to 448×448 pixels for ViT (Vision Transformer) compatibility.

Preprocessing Pipeline

- Resize Images
- Map Mask Values
- Validate Pairs
- Remove Unmatched

Total Classes

10 semantic categories:

- Background
- Trees
- Lush Bushes
- Dry Grass
- Dry Bushes
- Ground Clutter
- Logs
- Rocks
- Landscape
- Sky

Dataset: Offroad Segmentation Dataset

Model Architecture

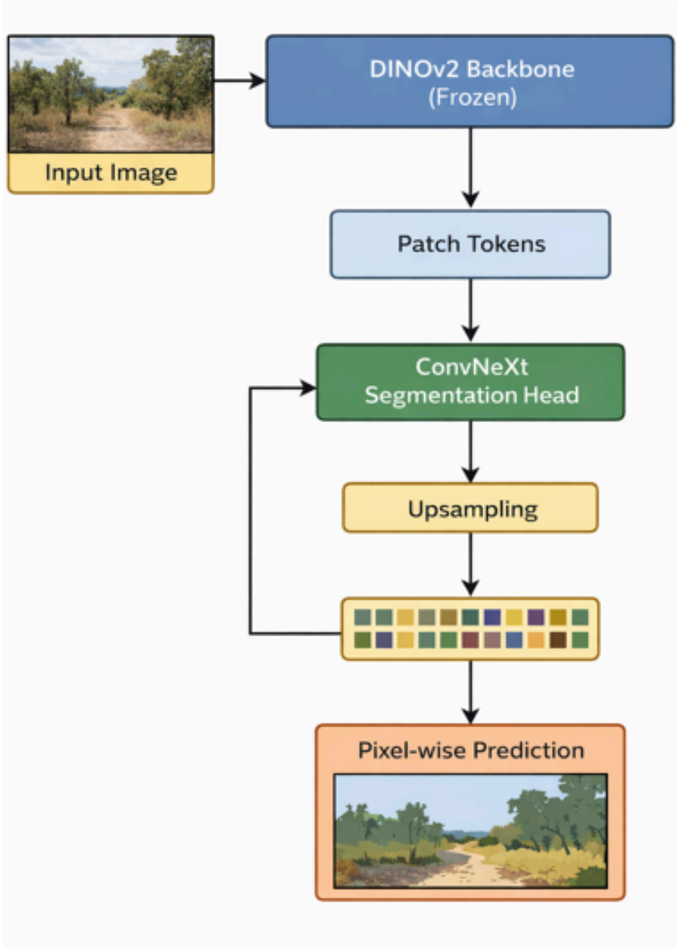

- DINOv2 Vision Transformer (ViT-S/14 backbone, self-supervised pretrained, frozen during initial training)
- ConvNeXt Segmentation Head (convolutional head reshaping patch tokens, 1×1 classifier)
- Upsampling & Output (bilinear upsampling for pixel-wise segmentation)

Training Setup

Hardware & Framework

GPU: NVIDIA RTX 3050 6GB
Framework: PyTorch 2.x
Python: 3.10
CUDA: 11.8

Hyperparameters

Batch Size: 2
Optimizer: AdamW
Learning Rate: Phase 1 1e-4 / Fine-tuning 5e-5
Scheduler: Cosine Annealing

Training Strategy

Phase 1 - Baseline Training

15 epochs with frozen backbone.

Initial Mean IoU: 0.4996

Phase 2 - Fine-Tuning

Additional 20 epochs, total ~35 epochs.

Reduced learning rate with cosine scheduling and mixed precision training.

Result

Significant improvement in segmentation consistency and IoU performance.

Final Results

Observed Improvements

- Better Boundary Prediction
- Small Object Detection
- Stable Classification

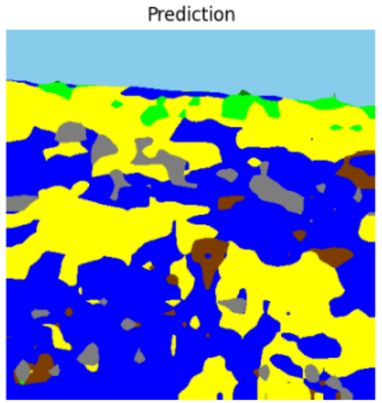

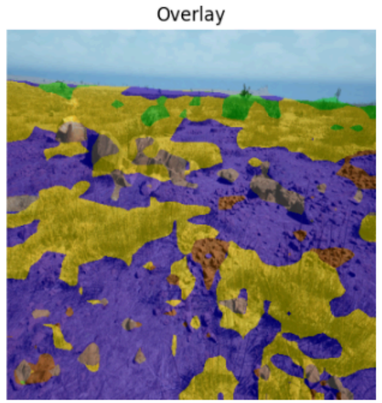

Key Metrics

- Final IoU (Improved from 0.4996 baseline) = Mean IoU: 0.4685958070134612 6

- Pixel Accuracy

- Total Epochs = 35

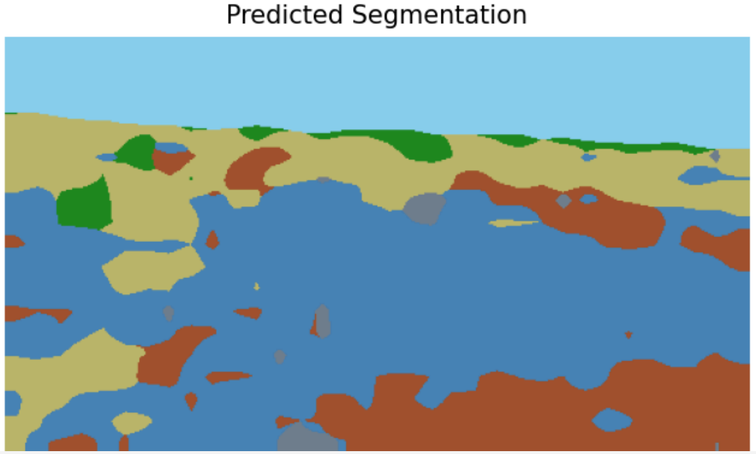

IoU / mAP Score

Intersection over Union (IoU = Intersection / Union)

Why IoU Matters

Standard metric for segmentation
Primary judging metric
Higher IoU indicates better accuracy
Accounts for precision and recall

Achievement

Achieved ~0.60+ IoU, a significant improvement from the 0.4996 baseline.

Challenges Faced

Addressing these key technical hurdles was crucial for the successful development and deployment of our off-road semantic segmentation model.

Dataset Path Configuration: Path errors and image-mask mismatch caused data loading issues.

Solution: Implemented automated filtering to ensure only valid image-mask pairs were used, preventing data corruption and training errors.

Environment Setup:

Ensuring perfect CUDA compatibility and dependency management was complex.

Solution: Established a clean Python 3.10 virtual environment with precise package versions to avoid conflicts and ensure stable operation.

GPU Memory Limitations

Challenge: Training was constrained by the 6GB VRAM of the NVIDIA RTX 3050.

Solution: Optimized memory usage by employing mixed precision training and a small batch size of 2, allowing larger models to fit.

Conclusion

Vision Transformers Excel
DINOv2's ViT architecture outperformed traditional CNNs at capturing global context in complex off-road environments.

Fine-tuning Improves IoU
Our two-phase training strategy, including fine-tuning, significantly boosted performance from a 0.4996 baseline to over 0.60+ IoU.

Proper Optimization is Crucial
Techniques like mixed precision training and hyperparameter tuning were essential for efficient and stable training within GPU memory limitations.

Thank You