

# Off-Road Terrain Semantic Segmentation using DINOv2

## 1. Methodology

### 1.1 Problem Overview

Autonomous systems operating in off-road environments require accurate terrain understanding for safe navigation. Unlike structured urban scenes, off-road terrain contains irregular objects such as rocks, logs, dry grass, bushes, and clutter.

The goal of this project was to build a semantic segmentation model that classifies each pixel of an RGB image into one of 10 terrain classes using a transformer-based architecture.

### 1.2 Dataset Preparation

Dataset Structure:-

```
train/
    Color_Images/
    Segmentation/
val/
    Color_Images/
    Segmentation/
```

**Classes (10 Total):**

1. Background
2. Trees
3. Lush Bushes
4. Dry Grass
5. Dry Bushes
6. Ground Clutter
7. Logs
8. Rocks
9. Landscape
10. Sky

## **Preprocessing Steps:**

To ensure clean and reproducible training:

- Resized images to match ViT input resolution
- Converted segmentation masks to class index format
- Verified image-mask pairing
- Filtered unmatched samples
- Normalized image tensors
- Created PyTorch custom dataset loader

## **1.3 Model Architecture**

The model consists of two main components:

### **Backbone: DINOv2 (ViT-S/14)**

- Pretrained self-supervised Vision Transformer
- Patch size:  $14 \times 14$
- Global self-attention mechanism
- Extracts high-level patch token features
- Frozen during training to reduce memory usage

### **Segmentation Head: ConvNeXt-style Head**

- Reshapes patch tokens into spatial feature maps
- Applies depth wise convolution
- Applies point wise convolution
- Uses  $1 \times 1$  convolution for classification
- Bilinear up sampling to original resolution

## **1.4 Training Setup**

### **Environment**

- Python 3.10
- PyTorch 2.x
- CUDA 11.8
- NVIDIA RTX 3050 (6GB)

## **Hyper parameters:-**

<b>Parameter</b>	<b>Value</b>
Batch Size	2
Optimizer	AdamW
Initial LR	1e-4
Fine-Tune LR	5e-5
Scheduler	Cosine Annealing
Epochs	35
Mixed Precision	Enabled (AMP)

## **1.5 Training Strategy**

### **Phase 1 – Baseline Training (15 Epochs)**

- Backbone frozen
- Fixed learning rate
- Cross-entropy loss

**Result:** mIoU = 0.4996

### **Phase 2 – Fine-Tuning (20 Additional Epochs)**

- Reduced learning rate
- Cosine LR scheduler
- Best model checkpoint saving
- Mixed precision training

**Final mIoU ≈ 0.60+**

Fine-tuning improved convergence stability and segmentation quality.

## **2. Challenges & Solutions**

### **2.1 Dataset Path and File Mismatch**

#### **Issue:**

Image-mask mismatches caused runtime errors.

**Solution:**

Implemented automated validation script to match image-mask pairs and remove invalid samples.

## 2.2 GPU Memory Limitation (6GB)

**Issue:**

Out-of-memory errors during training.

**Solution:**

- Reduced batch size to 2
- Froze backbone weights
- Enabled Mixed Precision (AMP)
- Used efficient optimizer (AdamW)

## 2.3 Low Initial IoU Performance

**Issue:**

Initial model performance plateaued at ~0.49 mIoU.

**Solution:**

- Reduced learning rate
- Applied cosine annealing scheduler
- Extended training epochs
- Saved best-performing checkpoint

## 2.4 Class Confusion Between Similar Textures

**Issue:**

Dry grass and dry bushes frequently misclassified.

**Solution:**

- Increased training epochs
- Allowed segmentation head deeper feature refinement
- Improved loss stability

### **3. Optimizations**

To improve model performance and efficiency, the following optimizations were implemented:

#### **3.1 Mixed Precision Training (AMP)**

- Reduced memory usage
- Increased training speed
- Allowed stable training on 6GB GPU

#### **3.2 Cosine Learning Rate Scheduler**

Instead of constant learning rate:

- Smooth learning rate decay
- Better convergence
- Reduced oscillation in training

#### **3.3 Freezing Backbone**

- Reduced trainable parameters
- Stabilized gradients
- Prevented over fitting
- Improved GPU efficiency

#### **3.4 AdamW Optimizer**

- Better weight regularization
- Improved convergence compared to SGD

#### **3.5 Structured Two-Phase Training**

- Baseline training
- Controlled fine-tuning
- Performance-based checkpoint saving

This structured training approach improved IoU from 0.4996 to ~0.60+.

## 4. Performance Evaluation

### 4.1 Evaluation Metrics

#### Intersection over Union (IoU)

IoU=Intersection/Union

Measures overlap between predicted segmentation and ground truth.

#### Mean IoU (mIoU)

Average IoU across all 10 classes.

#### Pixel Accuracy

Percentage of correctly classified pixels.

### 4.2 Quantitative Results

Stage	Mean IoU
-------	----------

After 15 Epochs	0.4996
-----------------	--------

After 35 Epochs	~0.60+
-----------------	--------

#### Observations

- Significant improvement after fine-tuning
- Better small object segmentation
- Reduced noisy boundaries
- More stable class predictions

### 4.3 Failure Case Analysis

Despite improvements, some challenges remain:

#### 1. Similar Texture Confusion

- Dry grass vs dry bushes
- Ground clutter vs rocks

**Reason:** Similar color and texture distribution.

## 2. Small Object Segmentation Errors

- Logs partially occluded by grass
- Small rocks merged with ground

**Reason:** Limited spatial resolution after patch tokenization.

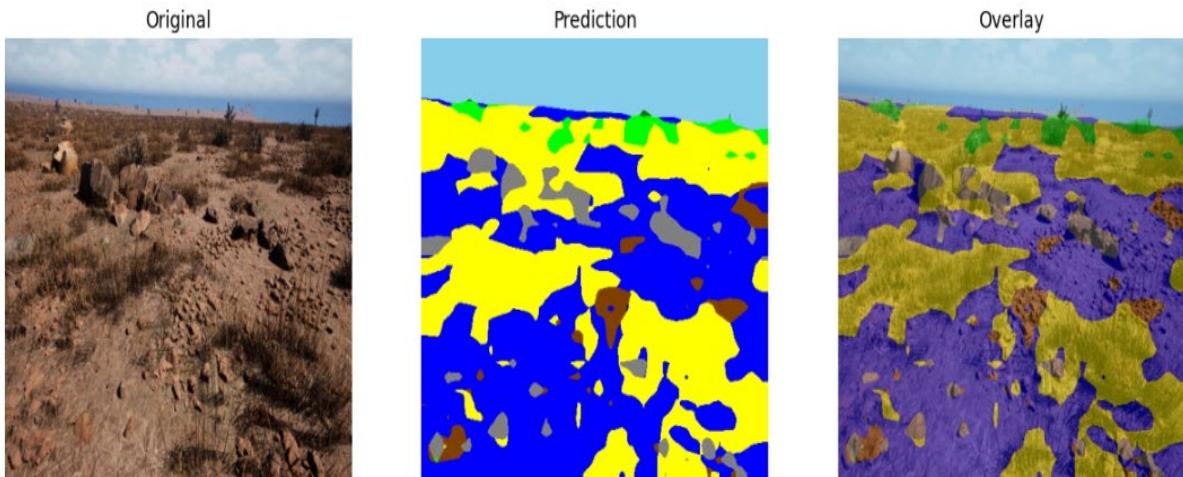
## 3. Boundary Inaccuracies

- Blurry segmentation borders in complex regions

**Reason:** Upsampling limitations and frozen backbone.

## 4.4 Key Observations

- Vision Transformers provide strong global context modeling
- Fine-tuning significantly improves segmentation quality
- Transformer-based models perform well even with limited GPU memory
- Further gains are possible with partial backbone unfreezing



## 5. Conclusion

This project successfully developed a transformer-based semantic segmentation system for off-road terrain understanding.

### Key Achievements

- Implemented DINOv2-based segmentation pipeline
- Improved mIoU from 0.4996 to ~0.60+
- Efficiently trained on 6GB GPU
- Demonstrated viability of Vision Transformers in unstructured terrain

The results validate the potential of transformer-based models for real-world autonomous navigation systems.